


SUTRA

The **AI Operating System** for Companies

AI does the work. Humans approve the outcomes.

THE BET

AI can already do the work.
Companies still **can't hand it over.**




88%
of companies use AI somewhere. Almost none let it run anything alone.

>40%
of agent projects scrapped by 2027 — the AI worked, the trust didn't (Gartner)

3
companies already run on Sutra — the wall, removed, in production

The Future Autonomous Company



One person can run a department. A few can run a company.

Humans sign. Fleets execute. **Nothing counts until it is checked and written down.**

Monday morning, one screen.

Neo · Mumbai ▾

THE ONE SCREEN YOU OPEN

Monday, 3 August 2026 · 08:04

NJ

TOP 3 ATTENTIONS — RANKED, WITH WHY IT REACHED YOU ON EVERY CARD

Custody vendor renewal — ₹3.4 Cr

above the finance charter's spend limit, so it stopped and came to you · alternatives priced, recommendation attached

SIGN

Q4 acquisition channel

4 experiments ran for 6 weeks · one clearly won · picking releases ₹80 L of budget to the winner

PICK

Compliance charter expires Friday

renewal drafted from this year's record · two authority limits proposed looser, one tighter

REVIEW

Start typing to act — "move the pitch", "why is CAC up?", "pause the sales fleet"...

GO

everything else — 18,400 jobs this month — ran under charter, on the record, without you

MY DAY — THE PLANNED DAY, DRAWN FROM YOUR TIME PRODUCT

09:30

Family-office pitch

brief ready

11:00

Board update

deck drafted

14:00

SEBI filing review

checked twice

16:30

Month-close

reconciled

FOCUS AREAS · THIS QUARTER

HNI segment build-out · week 6 of 12

Cost-to-income under 48% · on track

Two new fund partners · one signed

protected: 2 hrs deep work daily · the fleets schedule around it

Three calls need Nitin this morning — each arrives decided-ready. Eighteen thousand jobs this month did not need him, **and every one of them is on the record.**

CONFIDENTIAL

04 / 21

THE RULE-BOOK

The screen shows what stopped.
The charter is why.

Charter · Finance & Wealth Ops

OVERVIEW

DETAILS & SETTINGS

LIVE

owner: Nitin

GOALS — WHAT I ASKED FOR

Books closed by T+2, every month

✓ on track — closed T+1 last month

Zero payout errors

✓ 0 errors in 2,140 runs

Custody cost under 12 bps

✓ 10.8 bps now

GUARDRAILS I SET — IT NEVER CROSSES THESE

Never spend above ₹3 Cr without me

Never price below 60% margin

Client money — never touched without four eyes

9 additional guardrails active

RIGHT NOW

Custody renewal — ₹3.4 Cr — stopped at your ₹3 Cr line

waiting on you since 08:04

REVIEW

4 workflows running — month-close is at payout matching

✓ 18 asks answered this week · nothing overdue

last check 40s ago

Limits, tests, versions, the kill switch — one tap away in Details & settings. The engine enforces this page on every step.

A charter reads the way the owner thinks: **what I want, what it must never do, what is happening under it right now.** The enforcement machinery exists — it just lives a tap deeper.

CONFIDENTIAL

05 / 21

THE DIVISION OF LABOUR

Every company works in four places at once.
The machine can only hold one of them — so far.

AWARENESS ↑

KNOWN UNKNOWNNS

we know what we don't know

You know the question. The answer only shows up when you try something.

e.g. will tighter payment terms speed up collection? Does a different follow-up win more renewals?

MISSING: THE ANSWER

SUTRA PROBES IT — tries a controlled change, measures what happened, keeps what worked.

UNKNOWN UNKNOWNNS

we don't know what we don't know

Nobody asked. You find out when it hits you.

e.g. a rule change makes a service line illegal. Your biggest client's parent is acquired.

MISSING: THE QUESTION

SUTRA ANTICIPATES IT — standing watch on the weak signals, so you see it before it lands.

KNOWN KNOWNS

we know it, we do it by hand

You already know the process. A person still runs it, every month.

e.g. the month-end close, the payroll run, chasing overdue invoices, onboarding a client.

MISSING: OPERATIONALISATION

SUTRA OPERATIONALISES IT — the same work, run end to end and on the record.

UNKNOWN KNOWNS

we have it, we haven't worked it out

The pattern is already in your records. Nobody has found it.

e.g. which service lines actually make money after delivery cost. Which clients always pay late.

MISSING: THE WORKFLOW


SUTRA FINDS THE PATTERN — reads the trail in your own records, then makes it a workflow.

UNDERSTANDING →

Whatever you learn moves up and to the right. **What is settled, Sutra runs** — and the people you free up move outward, into judgement and the next question.

THE MISSING LAYER

Every layer of this stack has an owner. One doesn't.



The labs are climbing up from the model. The apps are reaching down from the work.

The band in between — **who is allowed to decide what, on whose authority, checked how, remembered where** — has no owner. It is also the only band that is different in every single company.

51%
of the \$37B enterprise AI spend already lands at the application end, not infrastructure

\$37B
enterprise AI spend in 2025 — 3.2× in a single year

— THE QUESTION EVERYONE ASKS

"Why won't Anthropic just build this?"

THREE REASONS

They say so themselves. The Claude Agent SDK exists so developers can build agents *on top of* Claude Code. And they gave away the plumbing — MCP, open-sourced, then donated to the Linux Foundation alongside OpenAI.

They contractually can't hold it. "By default, we will not use your inputs or outputs to train our models." The asset this layer runs on is *your* decision history — and they promised 300,000 customers they wouldn't accumulate it.

When a lab tries, it retreats. OpenAI shipped Agent Builder to own the enterprise agent layer in October 2025. It shuts down this November.

HONEST They are already partway here. The risk isn't that they build something better — it's that they build something good enough and bundle it free.

THE PRECEDENT, MEASURED

AWS has bundled CloudWatch since 2009.
Datadog still did \$3.43B last year, growing 28%.


AWS sells Redshift.
Snowflake still did \$2B through AWS's own marketplace.

Microsoft gives Copilot away inside seats people already pay for.
Glean is at ~\$300M ARR, growing 89%.

Bundling doesn't kill a layer whose value is per-company depth.

— THE MOAT

The system you install is not the one you run in six months — **it learned your company.**



changes need approval · reversible · all on the record


It watches its own misses and rewrites its own playbooks — with the owner's approval, reversible, recorded.

AI models are rented; everyone has the same one. **What the system learns about running your kind of company is owned** — and it compounds with every company it runs.

Said positively: **a horizontal vendor cannot hold this for every customer. It is per-company by definition.**

WHAT IT TAKES TO BUILD ONE

Autonomy is not installed.
It is **built, one loop at a time.**



SEE IT

Every decision lands in the record automatically — so you automate the real process, not the imagined one.

- you can query what happened

SHAPE IT

Frame first, then break the work down — each piece carrying the check that will prove it, written before work starts.

- every step has a done-test

RUN IT

Work executes and verifies itself. Nothing counts as finished until its own check passes; gates pause only for your calls.

- audit trail and a tested rollback

KEEP IT


The shape becomes a charter with an owner, a trigger and a switch-off rule — and the record makes the next loop better.

- owner, trigger, kill switch

Skip the first and you automate a process nobody runs. Skip the second and agents report "done" with no proof.
Skip the fourth and you are the scheduler, forever.


— HOW SUTRA IS BUILT

One flow. Work comes in, gets decided, gets done, gets kept.



MARKET

Bain sizes the work between the systems at \$100B. Almost none of it is taken.



\$13T

the US wage bill — against \$1.23T of software spend worldwide. AI competes with the first number.

6% → 2%

expected headcount growth, 2025 to 2026 — while 75% of CFOs raise technology budgets (Gartner)

16%

of enterprise AI deployments are true agents. The budget arrived; the autonomy didn't.

Bain's own framing is the pricing thesis, word for word: **converting labour cost into software spend**. The substitution is already visible in CFO budgets — that is the market opening, and it is opening now.

— SUTRA VS THE FIELD

Each of them has one or two.
The **intersection is empty.**

CAPABILITY	MICROSOFT Copilot	DUST	SIERRA	PALANTIR	ATLAN	SUTRA
Covers the whole company	yes	yes	one function	with heavy setup	data only	yes
Checks the work before it counts	no	watches after	yes	policy only	no	yes
Authority that grows as trust is earned	fixed roles	fixed roles	fixed	fixed roles	no	yes
Remembers decisions, not just documents	no	documents	scripts	data lineage	metadata	yes
Works above any AI model	own stack	yes	own stack	partial	yes	yes
Adapts itself per company	no	you configure	you configure	consultants	assisted	yes

Honest reading: **Dust owns coverage, Sierra owns verification** — both real, both well funded. Nobody has combined them, and nobody is attempting earned authority or a decision memory. That intersection is the company.

— ECONOMICS

Take one department.
The monthly close and payout reconciliation.

RUN BY A POD TODAY

₹30–90 L

a year, fully loaded

three to five people, their seats,
their tools and their supervisor



RUN ON SUTRA

₹20–30 L

a year, by function and volume

the same close, checked and on the record,
with a human signing what matters

Same workflow, two ways to run it — this is the cost of the *function*, not the cost of one employee.

This is not a cheaper-than-a-person argument. One Mumbai analyst costs a fraction of that pod, and less than we do. We are cheaper than the function, and cheaper than the heads you would add next year.

And the pod does not stay. Frontline attrition in Indian financial services runs from roughly half the team to the whole team, every year. The process knowledge leaves with them. Software does not resign.

— THE BET

YEAR TWO

\$10M

revenue · 125 customers × 3 functions

~\$140M

company worth, at 14x

YEAR FIVE

\$135M

revenue · 1,000 customers × 5 functions

~\$2.6B

company worth, at 19.5x

Valuation uses the nearest governance comparable early, then a high-growth software multiple at scale —
deliberately excluding the AI outlier premiums.

Price never moves: \$27k per function, all five years. This is the 80th percentile — roughly twice the classic SaaS top-decile path, about half Bessemer's breakout AI cohort.

The entry cost is five new customers a month, sustained. That needs a real sales team from month nine. Today: three companies run on it, and revenue is already not zero.

APPENDIX

Inside the product.

The screens the operator lives in · the platform underneath them · how the system grows.

One sentence holds it all: **the company runs; the person decides.**

8 OPERATING AREAS

58 CAPABILITIES ON THE BINDING LIST

5 GOVERNANCE GATES

Two conversations. Both running. You steer; nothing waits for you.

Chats · 2 running

SIMULTANEOUS · EACH UNDER ITS OWN CHARTER · BOTH ON THE RECORD

12:07

● Month-close · Finance fleet

RUNNING

✓ 11:48 · bank reconciliation — 6 of 6 accounts

✓ 11:56 · GST return filed · proof: ARN on file

▶ 12:07 · payout matching — 412 of 460...

NITIN · 12:02

hold vendor payouts till Thursday

FINANCE AGENT · 12:02

Held — 14 payouts moved to Thu 09:00. The close still lands Tuesday. Logged as your decision, reversible till Wednesday 18:00.

and release the two flagged ones|

SEND

you are typing here — the other agent keeps moving →

● HNI fee schedule · Sales fleet

WORKING NOW

✓ 12:04 · drafted 3 options · margins 61–74%

✓ 12:05 · margin-floor check — all pass

✓ 12:06 · RM feedback folded in — 2 of 2

▶ 12:07 · packaging option B with the compare...

SALES AGENT · 12:04

Option B is ready — recovers an estimated 2 of 3 lost mandates at –90bps of margin. Sign, or see the compare?

waiting on you · work continues on the parts that don't need you

SIGN

last tick 12:07 — it did not stop while you typed on the left

Work does not queue behind your attention. **Each agent runs its own loop under its own charter** — you drop into either conversation; the other never stops.

The platform inside: where work is given its rules.

Builder · Client onboarding — HNI

CHARTER: WEALTH OPS spend limit ₹50 L · kill switch ON

WORKFLOW LIBRARY

Standing24

Drafts3

Industry imports12

Retired7

THIS WORKFLOW

runs 340x/yr

94% without a human

₹41 per run

TALKS TO

CRM · custody API

KYC registry · email

EVERY STEP SAYS WHAT PROVES IT WORKED — AND WHETHER IT MAY RUN ALONE

STEP 1 · RUNS ALONE

KYC + risk profile

done when: registry check returns clean · profile scored

STEP 2 · RUNS ALONE

Fee schedule

done when: inside the margin floor · matches segment policy

STEP 3 · ASKS FIRST

Custody account

pauses: for a signature above ₹5 Cr initial funding

STEP 4 · RUNS ALONE

Welcome + first review

done when: client confirmed · review on both calendars

TEST RUN — against last month's real cases


PUBLISH UNDER CHARTER

publishing = versioned, reversible, on the record


Three ways a workflow is born: a person drafts it here · an industry template is imported and adapted · the system proposes it — from patterns it found in its own record.

One takeaway: **this is governed automation, built in the open** — the limits, the proofs and the pauses are set on this screen, not buried in code.

Five layers. Each consumes the one below. The screen is thin; the platform is deep.




Five systems, one loop.
The engine is the only mover — and it only borrows authority.



Read it as a sentence: **builders publish workflows → the engine runs them under charter authority → everything lands in the record → the record teaches the builders.** One loop — and every arrow is a screen in the product.

Day one it is a careful assistant. The record grows it into an operating system.



Both arms are driven by the same asset — the record — and gated by the same rule: **nothing gets more autonomy, and nothing gets hardened, without the run history to justify it.**